


Deep Learning


Andrew Ng

Thanks to Adam Coates, Kai Yu, Tong Zhang, Sameep Tandon,
Swati Dube, Brody Huval, Tao Wang,


Virtuous circle of AI



Data and machine learning



Virtuous circle of AI



Deep Learning



Adam Coates, Yoshua Bengio, Tom Dean, Jeff Dean, Nando de Freitas, Jeff Hawkins, Geoff Hinton, Quoc Le, Yann LeCun, Honglak Lee, Tommy Poggio, Ruslan Salakhutdinov, Yoram Singer, Josh Tenenbaum, Kai Yu, Tong Zhang,


Things we want to do with data

Images



Label image

Audio



Speech recognition

Text



Web search


STanford AI Robot (STAIR)



Computer vision: Identify coffee mug



Computer vision: Identify coffee mug



Why is computer vision hard?



The camera sees :


194	210	201	212	199	213	215	195	178	158	182	209
180	189	190	221	209	205	191	167	147	115	129	163
114	126	140	188	176	165	152	140	170	106	78	88
87	103	115	154	143	142	149	153	173	101	57	57
102	112	106	131	122	138	152	147	128	84	58	66
94	95	79	104	105	124	129	113	107	87	69	67
68	71	69	98	89	92	98	95	89	88	76	67
41	56	68	99	63	45	60	82	58	76	75	65
20	43	69	75	56	41	51	73	55	70	63	44
50	50	57	69	75	75	73	74	53	68	59	37
72	59	53	66	84	92	84	74	57	72	63	42
67	61	58	65	75	78	76	73	59	75	69	50

Computer vision




Learning
algorithm


Computer vision




Features for vision



SIFT






GIST





Shape context

Features for machine learning




Why is speech recognition hard?


Microphone recording:




Features for audio



Spectrogram




MFCC




Flux

Features for text



<DOC>
<DOCID> wsj94_008.0212 </DOCID>
<DOCNO> 940413-0062. </DOCNO>
<HL> Who's News:
<E_Burns Fry Ltd. </HL>
<DD> 04/13/94 </DD>
<SO> WALL STREET JOURNAL (J), PAGE B10 </SO>
<CO> MER </CO>
<IN> SECURITIES (SCR) </IN>
<TXT>
<p> BURNS FRY Ltd. (Toronto) -- Donald Wright, 4
named executive vice president and director of
brokerage firm. Mr. Wright resigned as president
Canada Inc., a unit of Merrill Lynch & Co., to
Kassirek, 48, who left Burns Fry last month. A
spokeswoman said it hasn't named a successor to
expected to begin his new position by the end of
</p>
</TXT>
</DOC>



Parser

Named entity

Stemming

The idea:

Most perception (input processing) in the brain may be due to one learning algorithm.




The idea:

Build learning algorithms
that mimic the brain.

Most of human intelligence may
be due to one learning algorithm.




The “one learning algorithm” hypothesis



[Roe et al., 1992]


The “one learning algorithm” hypothesis




[Metin & Frost, 1989]

Andrew Ng

Neurons in the brain



Neural Network (Deep Learning)



Deep Learning trends

Now



0-2 years

Tagged data

3-5 years

Tagged & untagged data



Learning from tagged data (supervised)



Coffee mug



Coffee mug



Coffee mug



Coffee mug



Coffee mug




Coffee mug



Testing: What is this?

Bigger is better



Google Brain




AI as a computer systems problem

10 million connections



1 billion connections

Speech recognition, and more....



[with Vincent Vanhoucke]

Deep Learning applications



Speech recognition




Image Search



Ads; Web search

Tagged vs. untagged data



Coffee mug



Coffee mug



Coffee mug



Coffee mug



Coffee mug




Coffee mug

Untagged data (unsupervised learning)



Unknown



Unknown



Unknown



Unknown




Unknown




Unknown

How does the brain process images?

Visual cortex looks for lines/edges.




Neuron #1 of visual cortex
(model)



Neuron #2 of visual cortex
(model)

Start with Image patches



152	147	128	84
129	113	107	87
98	95	89	88
60	82	58	76
51	73	55	70

$R^{14 \times 14}$

Sparse Coding

Input: Images patches $x^{(1)}, x^{(2)}, \dots$ (each in $\mathbb{R}^{14 \times 14}$)


Learn: Set of matrices $\phi_1, \phi_2, \dots, \phi_{64}$ (also $\mathbb{R}^{14 \times 14}$), so that each input x can be approximately written as a weighted sum of the ϕ_j 's:

$$x \approx \sum_{j=1}^{64} a_j \phi_j$$


s.t. a_j 's are mostly zero ("sparse")

Sparse Coding




$\phi_1, \phi_2, \dots, \phi_{64}$



Test example




$$x \approx 0.8 * \phi_{36} + 0.3 * \phi_{42} + 0.5 * \phi_{63}$$




Comparing to Biology


Brain (visual cortex)




Learning algorithm



Comparing to Biology



Comparing to Biology



Learning from YouTube videos



Unknown



Unknown



Unknown



Unknown




Unknown



Unknown

Face neuron




[Le et al., 2012]

Cat neuron



[Le et al., 2012]

Deep Learning



16,000 CPUs is expensive



GPUs (Graphics Processor Unit)



[Adam Coates, Bryan Catanzaro, et al.]

Building huge neural networks

10 million connections




1 billion connections




10 billion connections


Learning from tagged data



Highway perception




Deep Learning trends



Untagged data and AI (unsupervised learning)




Computer vision (~6 years)



Deep Learning

Speech recognition (next 2-3 years?)



Deep Learning

AI will transform the internet



Speech



Images



Text



Speech




Images

Speech recognition



Baidu Cool Box







35.5%




15.8%






Image queries



Clothing




Bags



Fruits & Vegetables

Baidu Eye





Andrew Ng

“Smart glasses” designs



Extending human perception



Extending human perception



Comparison to “smart glasses” designs



Andrew Ng



Speech



Images



Text

From Control to Perception



Andrew Ng

Stanford's PR-1 robot



[Ken Salisbury]


AI will transform the internet

Technology areas with potential for paradigm shift:

- Computer vision
- Speech recognition & speech synthesis
- Language understanding: Machine translation;
Web search; Dialog systems;
- Advertising
- Personalization/recommendation systems
- Robotics

All this is hard: scalability, algorithms.

Virtuous circle of AI



The AI mission





Thank you.


Thanks to Adam Coates, Yu Kai, Zhang Tong, Sameep Tandon,
Swati Dube, Brody Huval, Tao Wang,

Tutorial: deeplearning.stanford.edu


END

END

Discussion: Engineering vs. Data



Discussion: Engineering vs. Data





● Correctly found mug

● Mistake



● Correctly found mug

● Mistake



● Correctly found mug

● Mistake

